
Double-click (or enter) to edit

Download the `sampladata_orders.csv` dataset. As the name says it is a .csv file. The file contains information regarding sales order details of United Shipping Agency. Open it in a text editor or excel and inspect it first. Write a python program to load the data stored in a csv data file `sampladata_orders.csv` in a pandas DataFrame and perform the following operations.

- i. Print the top 5 row and top 10 row
- ii. Print the all column names
- iii. Print the all row index ranges
- iv. Print the Order Quantity column (top 5 values)
- v. Print The Sales column (top 10 values)
- vi. Print the row with ID:50
- vii. Print the third row
- viii. Print the Order Quantity, Sales, Discount and Profit of the 2nd,4th, 6th and 8th row:
- ix. Print the Order Quantity, Sales, Discount and Profit of orders with discount > 10%
- x. Print the production cost.(production cost= Sales-Profit)

```
import pandas as pd
orders = pd.read_csv("/content/sampladata_orders.csv")
#print("First 5 entries:")
#print(orders.head())
#print(orders.head(10))
#print(orders.columns) #Print the all column names
#Print the all row index ranges
#print(orders.index)
#print(orders["Order Quantity"].head()) #Print the Order Quantity column (top 5 values)
#print(orders["Sales"].head(10))#Print The Sales column (top 10 values)
#print("The row with ID:50")
#r50 = orders.loc[50] #Print the row with ID:50
#print(r50)
#r3=orders.iloc[3]
#print(r3) #Print the third row
#print(orders[2:9:2][["Order Quantity", "Sales","Discount", "Profit"]])
#print(orders[orders["Discount"] > 0.1][["Order Quantity","Sales","Discount", "Profit"]])
pcost=orders["Sales"]-orders["Profit"]
print(pcost)
```

```
0      474.7900
1     9665.2100
2      197.8600
3     3766.7895
4      363.3300
```

```
      ...  
8394    1617.2200  
8395     370.3200  
8396     480.7300  
8397     629.0775  
8398    1445.5800  
Length: 8399, dtype: float64
```

Download the diamonds.csv dataset. As the name says it is a .csv file. The file contains information regarding diamonds details of Laisha diamond shop. Open it in a text editor or excel and inspect it first. Write a python program to load the data stored in a csv file diamonds.csv in a pandas DataFrame and perform the following operations.

- I. find carat weight at least 0.3 and write this detail into cart.csv file
- II. find total price of carat weight at maximum 0.3
- III. Sort the 'carat' column by ascending and descending order.
- IV. Print the premium diamond details
- V. Find the most expensive diamond whose color is D

```
import pandas as pd  
df = pd.read_csv('/content/diamonds.csv')  
sample=df[df["carat"]<0.3]  
sample.to_csv('cart.csv')  
tot=df[df["carat"]>0.3]  
tot["price"].sum()  
result = df.sort_values('carat')  
print(result)  
result = df.sort_values('carat', ascending=False)  
print(result)  
pre=df.groupby("cut")  
pre.get_group("Premium")
```

	carat	cut	color	clarity	depth	table	price	x	y	z
31593	0.20	Premium	E	VS2	61.1	59.0	367	3.81	3.78	2.32
31597	0.20	Ideal	D	VS2	61.5	57.0	367	3.81	3.77	2.33
31596	0.20	Premium	F	VS2	62.6	59.0	367	3.73	3.71	2.33
31595	0.20	Ideal	E	VS2	59.7	55.0	367	3.86	3.84	2.30
31594	0.20	Premium	E	VS2	59.7	62.0	367	3.84	3.80	2.28
...
25999	4.01	Premium	J	I1	62.5	62.0	15223	10.02	9.94	6.24
25998	4.01	Premium	I	I1	61.0	61.0	15223	10.14	10.10	6.17
27130	4.13	Fair	H	I1	64.8	61.0	17329	10.00	9.85	6.43
27630	4.50	Fair	J	I1	65.8	58.0	18531	10.23	10.16	6.72
27415	5.01	Fair	J	I1	65.5	59.0	18018	10.74	10.54	6.98

[53940 rows x 10 columns]

	carat	cut	color	clarity	depth	table	price	x	y	z
27415	5.01	Fair	J	I1	65.5	59.0	18018	10.74	10.54	6.98
27630	4.50	Fair	J	I1	65.8	58.0	18531	10.23	10.16	6.72
27130	4.13	Fair	H	I1	64.8	61.0	17329	10.00	9.85	6.43
25999	4.01	Premium	J	I1	62.5	62.0	15223	10.02	9.94	6.24
25998	4.01	Premium	I	I1	61.0	61.0	15223	10.14	10.10	6.17
...
31592	0.20	Premium	E	VS2	59.0	60.0	367	3.81	3.78	2.24
31591	0.20	Premium	E	VS2	59.8	62.0	367	3.79	3.77	2.26
31601	0.20	Premium	D	VS2	61.7	60.0	367	3.77	3.72	2.31
14	0.20	Premium	E	SI2	60.2	62.0	345	3.79	3.75	2.27
31596	0.20	Premium	F	VS2	62.6	59.0	367	3.73	3.71	2.33

[53940 rows x 10 columns]

	carat	cut	color	clarity	depth	table	price	x	y	z
1	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31
3	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63
12	0.22	Premium	F	SI1	60.4	61.0	342	3.88	3.84	2.33

1. WAP program to read a sampledata_orders.csv file

and perform descriptive statics and print the result.

```
import pandas as pd
orders = pd.read_csv("/content/sampledata_orders.csv")
print("A description of the numerical values:")
print(orders.describe())
orders.min()
orders.max()
orders.std()
orders.var()
orders.mean()
orders.median()
orders.mode()
orders.sum()
orders.count()
orders.rank()
```

```
orders.skew()
orders.kurt()
```

A description of the numerical values:

	Row ID	Order ID	Order Quantity	Sales	Discount \
count	8399.000000	8399.000000	8399.000000	8399.000000	8399.000000
mean	4200.000000	29965.179783	25.571735	1775.878179	0.049671
std	2424.726789	17260.883447	14.481071	3585.050525	0.031823
min	1.000000	3.000000	1.000000	2.240000	0.000000
25%	2100.500000	15011.500000	13.000000	143.195000	0.020000
50%	4200.000000	29857.000000	26.000000	449.420000	0.050000
75%	6299.500000	44596.000000	38.000000	1709.320000	0.080000
max	8399.000000	59973.000000	50.000000	89061.050000	0.250000

	Profit	Unit Price	Shipping Cost	Product Base Margin
count	8399.000000	8399.000000	8399.000000	8336.000000
mean	181.184423	89.346259	12.838557	0.512513
std	1196.653372	290.354383	17.264052	0.135589
min	-14140.700000	0.990000	0.490000	0.350000
25%	-83.315000	6.480000	3.300000	0.380000
50%	-1.500000	20.990000	6.070000	0.520000
75%	162.750000	85.990000	13.990000	0.590000
max	27220.690000	6783.020000	164.730000	0.850000

<ipython-input-7-ec6ac894d235>:7: FutureWarning: Dropping of nuisance columns in
orders.std()

<ipython-input-7-ec6ac894d235>:8: FutureWarning: Dropping of nuisance columns in
orders.var()

<ipython-input-7-ec6ac894d235>:9: FutureWarning: Dropping of nuisance columns in
orders.mean()

<ipython-input-7-ec6ac894d235>:10: FutureWarning: Dropping of nuisance columns
orders.median()

<ipython-input-7-ec6ac894d235>:15: FutureWarning: Dropping of nuisance columns
orders.skew()

<ipython-input-7-ec6ac894d235>:16: FutureWarning: Dropping of nuisance columns
orders.kurt()

Row ID -1.200000

Order ID -1.178317

Order Quantity -1.208020

Sales 60.928376

Discount -0.959411

Profit 67.349705

Unit Price 271.168733

Shipping Cost 7.751587

Product Base Margin -0.660870

dtype: float64

2. Write a python program to read a sampledata_orders.csv file perform descriptive statics for the first 25 rows of the sales column and print the result.

```
import pandas as pd
orders = pd.read_csv("sampledata_orders.csv")
```

```

orders[0:25:1][["Sales"]].min()
orders[0:25:1][["Sales"]].max()
orders[0:25:1][["Sales"]].std()
orders[0:25:1][["Sales"]].var()
orders[0:25:1][["Sales"]].mean()
orders[0:25:1][["Sales"]].median()
orders[0:25:1][["Sales"]].mode()
orders[0:25:1][["Sales"]].sum()
orders[0:25:1][["Sales"]].count()
orders[0:25:1][["Sales"]].rank()
orders[0:25:1][["Sales"]].skew()
orders[0:25:1][["Sales"]].kurt()


```

Horizontal bar chart

```

import numpy as np
import pandas as pd
#!pip install matplotlib==3.4
from matplotlib import pyplot as plt
plt.style.use('seaborn-whitegrid')
import pandas as pd
import numpy as np
fig = plt.figure()
ax = plt.axes()

```




```

from matplotlib import pyplot as plt
plt.style.use('seaborn-whitegrid')
fig = plt.figure()
ax = plt.axes()
data = {'Apple': 10,
        'Orange': 15,
        'Banana': 13,
        'Pear': 9,
        'Pomegrante': 11,
        'Lemon': 5,
        'Lime': 20}
names = list(data.keys())
values = list(data.values())
ax = plt.axes()
ax.barh(names, values, color='c', height=0.80)
ax.set_xlabel('Quantity')

```

Horizontal bar-chart

```
ax.set_ylabel('Fruit')
ax.set_title('Fruit Quantities');
```




```
labels = ['G1', 'G2', 'G3', 'G4', 'G5']
men_means = [20, 34, 30, 35, 27]
women_means = [25, 32, 34, 20, 25]
x = np.arange(len(labels)) # the label locations
width = 0.25 # the width of the bars
fig, ax = plt.subplots()
rects1 = ax.bar(x - width/2, men_means, width, label='Men')
rects2 = ax.bar(x + width/2, women_means, width, label='Women')
ax.set_ylabel('Scores')
ax.set_title('Scores by group and gender')
ax.set_xticks(x)
ax.set_xlabel('Groups')
ax.set_xticklabels(labels)
ax.legend();

# plt.show()
```



```
fig, ax = plt.subplots()
ax.grid(linestyle='--', color='0.75', axis = 'y')
ax.set_axisbelow(True) # Grid behind bars
p1 = ax.bar(labels, men_means, width, label='Men')
p2 = ax.bar(labels, women_means, width, bottom=men_means,
            label='Women')
ax.set_ylabel('Scores')
ax.set_title('Scores by group and gender')
ax.legend()
# Label with label_type 'center' instead of the default 'edge'
ax.bar_label(p1, label_type='center')
```

```
ax.bar_label(p2, label_type='center')
ax.bar_label(p2)
ax.set_ylim(0,70)
```



flights Data set

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
flights = pd.read_csv("/content/flights.csv")
flights.head()
```

	year	month	day	dep_time	dep_delay	arr_time	arr_delay	carrier	tailnum	1
0	2013	1	1	517.0	2.0	830.0	11.0	UA	N14228	
1	2013	1	1	533.0	4.0	850.0	20.0	UA	N24211	
2	2013	1	1	542.0	2.0	923.0	33.0	AA	N619AA	
3	2013	1	1	554.0	-6.0	812.0	-25.0	DL	N668DN	
4	2013	1	1	554.0	-4.0	740.0	12.0	UA	N39463	

```
flights.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 160754 entries, 0 to 160753
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   year        160754 non-null  int64
1   month       160754 non-null  int64
2   day         160754 non-null  int64
3   dep_time    158418 non-null  float64
4   dep_delay   158418 non-null  float64
5   arr_time    158275 non-null  float64
6   arr_delay   157927 non-null  float64
7   carrier     160754 non-null  object
8   tailnum     159321 non-null  object
9   flight      160754 non-null  int64
```

```

10 origin      160754 non-null object
11 dest        160754 non-null object
12 air_time    157927 non-null float64
13 distance    160754 non-null int64
14 hour        158418 non-null float64
15 minute      158418 non-null float64
dtypes: float64(7), int64(5), object(4)
memory usage: 19.6+ MB

```

```

# Select the rows that have at least one missing value
flights[flights.isnull().any(axis=1)]

```

	year	month	day	dep_time	dep_delay	arr_time	arr_delay	carrier	tailr
330	2013	1	1	1807.0	29.0	2251.0	NaN	UA	N314
403	2013	1	1	NaN	NaN	NaN	NaN	AA	N3EH
404	2013	1	1	NaN	NaN	NaN	NaN	AA	N3EV
855	2013	1	2	2145.0	16.0	NaN	NaN	UA	N122
858	2013	1	2	NaN	NaN	NaN	NaN	AA	N
...
159681	2013	9	28	1214.0	-11.0	1801.0	NaN	AA	N48E
159854	2013	9	28	NaN	NaN	NaN	NaN	AA	N320
159855	2013	9	28	NaN	NaN	NaN	NaN	US	N
160185	2013	9	29	1734.0	23.0	2159.0	NaN	UA	N463
160286	2013	9	29	NaN	NaN	NaN	NaN	UA	N

2827 rows × 16 columns

2. Filter all the rows where arr_delay value is missing:

```

flights1 = flights[ flights['arr_delay'].notnull( )]
flights1.head(10)

```


	year	month	day	dep_time	dep_delay	arr_time	arr_delay	carrier	tailnum	1
0	2013	1	1	517.0	2.0	830.0	11.0	UA	N14228	
1	2013	1	1	533.0	4.0	850.0	20.0	UA	N24211	
2	2013	1	1	542.0	2.0	923.0	33.0	AA	N619AA	
3	2013	1	1	554.0	-6.0	812.0	-25.0	DL	N668DN	
4	2013	1	1	554.0	-4.0	740.0	12.0	UA	N39463	

3. Remove all the observations with missing values

```
flights2 = flights.dropna()
flights2
```

	year	month	day	dep_time	dep_delay	arr_time	arr_delay	carrier	tailnum
0	2013	1	1	517.0	2.0	830.0	11.0	UA	N14228
1	2013	1	1	533.0	4.0	850.0	20.0	UA	N24211
2	2013	1	1	542.0	2.0	923.0	33.0	AA	N619AA
3	2013	1	1	554.0	-6.0	812.0	-25.0	DL	N668DN
4	2013	1	1	554.0	-4.0	740.0	12.0	UA	N39463
...
160749	2013	9	30	2105.0	-1.0	2329.0	-25.0	UA	N477
160750	2013	9	30	2121.0	21.0	2349.0	-25.0	DL	N193
160751	2013	9	30	2140.0	0.0	10.0	-30.0	AA	N335
160752	2013	9	30	2149.0	-7.0	2245.0	-23.0	UA	N813
160753	2013	9	30	2233.0	80.0	112.0	42.0	UA	N576

157927 rows × 10 columns

4. Fill missing values with zeros for dep_delay

```
nomiss = flights['dep_delay'].fillna(0)
nomiss
```

```
nomiss.isnull().any()
```

False

5. Let's compute summary statistic per a group':

6. Using `agg()` methods for aggregation fetch min,max and mean for columns `dep_delay` and `arr_delay`:


```
flights[['dep_delay', 'arr_delay']].agg(['min', 'mean', 'max'])
```

	dep_delay	arr_delay
min	-33.000000	-75.000000
mean	9.463773	2.094537
max	1014.000000	1007.000000

7. Using seaborn plot factorplot

```
import seaborn as sns
sns.catplot(x='carrier', y='dep_delay', data=flights, kind='bar')
```

```
<seaborn.axisgrid.FacetGrid at 0x7fbdeb0ed760>
```



Violin plot

```
sns.violinplot(x = "dep_delay", data=flights)
```

```
<Axes: xlabel='dep_delay'>
```



```
sns.scatterplot(x='carrier', y='dep_delay', data=flights)
```

```
sns.jointplot(x='carrier', y='dep_delay', data=flights)
```

```
sns.boxplot(x='carrier',y='dep_delay', data=flights)
```

```
sns.swarmplot(x='carrier',y='dep_delay', data=flights)
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-18-85e526d2b4d0> in <module>
----> 1 sns.swarmplot(x='carrier',y='dep_delay', data=flights)
```

14 frames

```
/usr/local/lib/python3.9/dist-packages/seaborn/categorical.py in
first_non_overlapping_candidate(self, candidates, neighbors)
    3502             dx = neighbors_x - x_i
    3503             dy = neighbors_y - y_i
-> 3504             sq_distances = np.square(dx) + np.square(dy)
    3505
    3506             sep_needed = np.square(neighbors_r + r_i)
```

KeyboardInterrupt:

SEARCH STACK OVERFLOW

Error in callback <function flush_figures at 0x7fbe1ea9f0d0> (for post_execute)

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
/usr/local/lib/python3.9/dist-packages/ipykernel/pylab/backend_inline.py in
flush_figures()
    119         # ignore the tracking, just draw and close all figures
    120         try:
-> 121             return show(True)
    122         except Exception as e:
    123             # safely show traceback if in IPython, else raise
```

18 frames

```
<decorator-gen-2> in __call__(self, obj)

/usr/local/lib/python3.9/dist-packages/seaborn/categorical.py in
first_non_overlapping_candidate(self, candidates, neighbors)
    3509         # squared distance between candidate and any of the
neighbors has
    3510         # to be at least square of the summed radii
-> 3511         good_candidate = np.all(sq_distances >= sep_needed)
    3512
    3513         if good_candidate:
```

KeyboardInterrupt:

SEARCH STACK OVERFLOW

sns.pairplot(flights)

Salaries data set

```
import pandas as pd
df = pd.read_csv("/content/Salaries.csv")
```

```
df.head(10)
```

Rename the columns

```
df_new =df.rename(columns={
    'discipline': 'subject',
    'sex': 'gender'
})
df_new
```

Create a new column

```
df = df.assign( salary_k = lambda x: x.salary/1000.0)
df.head(10)
```

	rank	discipline	phd	service	sex	salary	salary_k
0	Prof	B	56	49	Male	186960	186.960
1	Prof	A	12	6	Male	93000	93.000
2	Prof	A	23	20	Male	110515	110.515
3	Prof	A	40	31	Male	131205	131.205
4	Prof	B	20	18	Male	104800	104.800
5	Prof	A	20	20	Male	122400	122.400
6	AssocProf	A	20	17	Male	81285	81.285
7	Prof	A	18	18	Male	126300	126.300
8	Prof	A	29	19	Male	94350	94.350
9	Prof	A	51	51	Male	57800	57.800

Check how many unique values in a column

```
df['rank'].unique()

array(['Prof', 'AssocProf', 'AsstProf'], dtype=object)
```

Get frequency table for a categorical or binary column

```
df['rank'].value_counts()
```

Calculate the mean salary for men and women.

```
df.groupby('sex')['salary'].mean()
```


Group using 2 variables - sex and rank:

```
df.groupby(['rank','sex'], sort=True)[['salary']].mean()
```

```
%matplotlib inline
import matplotlib.pyplot as plt
plt.hist(df['salary'],bins=8, density=True)
```

```
df.groupby(['rank'])['salary'].count().plot(kind='bar')
```

<Axes: xlabel='rank'>



Data Cleaning

```
import pandas as pd
data =pd.read_csv('/content/data.csv')
print(data)
```

```
data.describe()
```

```
data.rank()
```

```
data.isnull()
```

```
data.dropna()
```

```
from numpy import NaN
data.replace({NaN:0.00})
```

```
data.fillna(3)
```

```
data.fillna(method='pad')
```

```
data.fillna(method='backfill') #Example of replacing missing values by filling backward
```

```
data.drop_duplicates()
```

```
del data['YOB']  
print(data)
```

```
print(data.rename(columns={'Name': 'FirstName', 'Surname': 'LastName'}))
```

Simple bar chart

```
data = {'Apple': 10, 'Orange': 15, 'Lemon': 5, 'Lime': 20}  
names = list(data.keys())  
values = list(data.values())
```

```
fig = plt.figure()  
ax = plt.axes()  
# ax.grid(linestyle='--', color='0.85') # Color: grayscale between 0 and 1
```

```
ax.bar(names, values);
```

```
data = {'Apple': 10, 'Orange': 15, 'Lemon': 5, 'Lime': 20, 'Bannan': 30, 'WM': 18, 'mango'  
names = list(data.keys())  
values = list(data.values())
```

```
fig = plt.figure(figsize=(10,7))  
ax = plt.axes()
```

```
ax.bar(names, values, color='r', width = 0.40);  
ax.set_xlabel('Fruit')  
ax.set_ylabel('Quantity')  
ax.set_title('Fruit Quantities');
```

```
data = {'Apple': 10,  
        'Orange': 15,  
        'Banana': 13,  
        'Pear': 9,  
        'Pomegrante': 11,  
        'Lemon': 5,  
        'Lime': 20}  
names = list(data.keys())  
values = list(data.values())
```

```
ax = plt.axes()
```

```
ax.barh(names, values, color='c', height=0.80) # Horizontal bar-chart
```

```

ax.set_xlabel('Quantity')
ax.set_ylabel('Fruit')
ax.set_title('Fruit Quantities');

labels = ['G1', 'G2', 'G3', 'G4', 'G5']
men_means = [20, 34, 30, 35, 27]
women_means = [25, 32, 34, 20, 25]

x = np.arange(len(labels)) # the label locations
width = 0.40 # the width of the bars

fig, ax = plt.subplots()
rects1 = ax.bar(x - width/2, men_means, width, label='Men')
rects2 = ax.bar(x + width/2, women_means, width, label='Women')

# # Add some text for labels, title and custom x-axis tick labels, etc.
ax.set_ylabel('Scores')
ax.set_xlabel('Groups')
ax.set_title('Scores by group and gender')
ax.set_xticks(x)
ax.set_xticklabels(labels)
ax.legend();
plt.show()

```

Scatter Plot

```

x = np.random.randn(20)
y = np.random.randn(20)
fig, ax = plt.subplots()
ax.scatter(x, y);
fig, ax = plt.subplots() # a figure with a single Axes
ax.scatter(x, y, marker = "v");
fig, axs = plt.subplots(2, 3, sharex=True, sharey=True, figsize=(16,12));

# marker symbol
axs[0, 0].scatter(x, y, s=80, marker=">")
axs[0, 0].set_title("Right Triangle ")
# marker from TeX
axs[0, 1].scatter(x, y, s=80, marker=r'$\alpha$')
axs[0, 1].set_title(r"marker=r'\$\alpha\$'")

# marker from path
verts = [[-1, -1], [1, -1], [1, 1], [-1, -1]]
axs[0, 2].scatter(x, y, s=80, marker=verts)
axs[0, 2].set_title("marker=verts")

# regular polygon marker
axs[1, 0].scatter(x, y, s=80, marker=(5, 0))
axs[1, 0].set_title("polygon")

# regular star marker
axs[1, 1].scatter(x, y, s=80, marker=(5, 1))
axs[1, 1].set_title("star")

# regular asterisk marker
axs[1, 2].scatter(x, y, s=80, marker=(5, 2))
axs[1, 2].set_title("asterisk");

```